

# Principles of Machine Learning

## Lecture 7: Model Evaluation and Validation

Sharif University of Technology  
Dept. of Aerospace Engineering

April 20, 2025



# Table of Contents

- 1 Introduction
- 2 Performance Measures
- 3 Error Analysis
- 4 Cross-validation & Bootstrap



# Table of Contents

- 1 Introduction
- 2 Performance Measures
- 3 Error Analysis
- 4 Cross-validation & Bootstrap



## Model Evaluation

- Some performance measures/metrics  $\rightarrow$  effectiveness of the model
- For regression tasks:
  - **Mean Absolute Error**
  - **Mean Squared Error**
  - **Root Mean Squared Error**
  - R-squared ( $R^2$ )
- For classification tasks:
  - ???
  - :
  - ???



## MNIST Dataset


Before continuing on evaluating classifiers, let's get to know about the MNIST dataset, the “Hello World!” of the AI/ML!

- Modified National Institute of Standards and Technology dataset was created in 1998
- A set of 70,000 small images of digits handwritten by high school students and employees of the US Census Bureau
- Each image is labeled with the digit it represents and has 784 features (pixels)



# Introduction

## MNIST Dataset



(a) MNIST sample belonging to the digit '7'.



(b) 100 samples from the MNIST training set.



# Table of Contents

- 1 Introduction
- 2 Performance Measures**
- 3 Error Analysis
- 4 Cross-validation & Bootstrap



## Accuracy

- The most basic performance measure
- Defined as *the ratio of correct predictions*

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{All Predictions}}$$





## Example 1. Implementing a classifier for MNIST

```
1 from sklearn.datasets import fetch_openml
2 mnist = fetch_openml('mnist_784', as_frame=False)
3
4 X, y = mnist.data, mnist.target
5 X_train, X_test, y_train, y_test = X[:60000], X[60000:], y
   [:60000], y[60000:]
6
7 y_train_5 = (y_train == '5') # True for all 5s, False for all
   other digits
8 y_test_5 = (y_test == '5')
```



## Example 1 contd. Implementing a classifier for MNIST

```
1 from sklearn.linear_model import SGDClassifier
2
3 sgd_clf = SGDClassifier(random_state=42)
4 sgd_clf.fit(X_train, y_train_5)
5
6 >>> sgd_clf.predict([some_digit])
7 array([ True])
8
```



# Performance Measures – Accuracy using Cross Validation

## Example 1 contd. Implementing a classifier for MNIST

```
1 >>> from sklearn.model_selection import cross_val_score
2 >>> cross_val_score(sgd_clf, X_train, y_train_5, cv=3, scoring=
    "accuracy")
3 array([0.95035, 0.96035, 0.9604 ])
4
5 from sklearn.dummy import DummyClassifier
6 dummy_clf = DummyClassifier()
7 dummy_clf.fit(X_train, y_train_5)
8 print(any(dummy_clf.predict(X_train))) # prints False: no 5s
    detected
9
10 >>> cross_val_score(dummy_clf, X_train, y_train_5, cv=3,
    scoring="accuracy")
11 array([0.90965, 0.90965, 0.90965])
12
```



## Cross Validation

- A model validation technique
- Splitting to training data and validation data
- Better generalization → reduce the chance of overfitting

## $k$ -fold Cross-validation

- Splitting the training set into  $k$  folds
- Training the model  $k$  times
- Holding out a different fold each time for evaluation



## Why not “accuracy”?

- From Example 1.  $\longrightarrow$  *accuracy* is usually not a good performance measure for classifiers
- Especially for *skewed* datasets (i.e., some classes are more frequent than others)



## Confusion Matrices

- General idea: count the number of times instances of class A are classified as class B, for all A/B pairs
- First need to have a set of predictions so that they can be compared to the actual targets
- Could make predictions on the test set, but it's best to keep that untouched for now
- Each row: an actual class
- Each column: a predicted class



# Performance Measures – Confusion Matrix

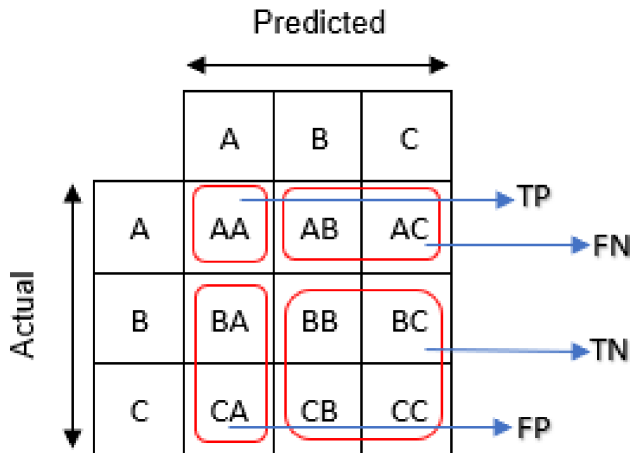
## A Confusion Matrix

|               |          | Predicted Values |          |
|---------------|----------|------------------|----------|
|               |          | Positive         | Negative |
| Actual Values | Positive | TP               | FN       |
|               | Negative | FP               | TN       |



# Performance Measures – Confusion Matrix

A  $3 \times 3$  Confusion Matrix example





# Performance Measures – Confusion Matrix

## Implementing CM for Example 1.

```
1 from sklearn.model_selection import cross_val_predict
2 y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5,
   cv=3)
3
4 >>> from sklearn.metrics import confusion_matrix
5 >>> cm = confusion_matrix(y_train_5, y_train_pred)
6 >>> cm
7 array([[53892, 687],
8        [ 1891, 3530]])
9
10 >>> y_train_perfect_predictions = y_train_5 # pretend we
   reached perfection
11 >>> confusion_matrix(y_train_5, y_train_perfect_predictions)
12 array([[54579, 0],
13        [ 0, 5421]])
14
```



# Performance Measures – Confusion Matrix

The confusion matrix gives useful metrics:

- **Precision:** the accuracy of the positive predictions

$$\text{Precision} = \frac{TP}{TP + FP}$$

- **Recall**, sensitivity or the true positive rate (TPR): the ratio of positive instances that are correctly detected by the classifier


$$\text{Recall} = \frac{TP}{TP + FN}$$

- TP: # true positives
- FP: # false positives
- TN: # true negatives
- FN: # false negatives



# Performance Measures – Confusion Matrix

The confusion matrix for the example 1 (predicting the digit 5)



## Implementing Precision, Recall, & F-1 for Example 1.

```
1 >>> from sklearn.metrics import precision_score, recall_score
2 >>> precision_score(y_train_5, y_train_pred) # == 3530 / (687 +
      3530)
3 0.8370879772350012
4 >>> recall_score(y_train_5, y_train_pred) # == 3530 / (1891 +
      3530)
5 0.6511713705958311
6
7 >>> from sklearn.metrics import f1_score
8 >>> f1_score(y_train_5, y_train_pred)
9 0.7325171197343846
10
```



## F-1 Score

- A combination of precision and recall
- Defined as the harmonic mean of the precision and recall


$$F-1 = 2 \cdot \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

- Useful especially when comparing two classifiers
- Both precision and recall high  $\rightarrow$  high F-1 score



# Performance Measures – The Precision/Recall Trade-off

The precision/recall trade-off is carried out by a *decision function* that determines a *threshold*:



# Performance Measures – The Precision/Recall Trade-off





Figure: Precision and recall versus the decision threshold



# Performance Measures – The Precision/Recall Trade-off

Another way of selecting a good precision/recall trade-off:






## The receiver operating characteristic (ROC) curve

- A common evaluation tool for binary classifiers
- Plots TPR (true positive rate) against FPR (false positive rate)
- $\text{TPR} = \text{recall: } \textit{sensitivity}$ ,  $\text{FPR} = 1 - \text{TNR: } 1 - \textit{specificity}$



# Performance Measures – The Precision/Recall Trade-off

## The ROC curve




## For **The ROC curve**

- A good classifier: stay toward the top-left corner
- To compare two classifiers: measure the *area under the curve* (AUC)
  - Perfect:  $AUC=1$
  - A purely random classifier:  $AUC=0.5$



# Performance Measures – The Precision/Recall Trade-off



**Figure:** Comparing PR curves: the random forest classifier is superior to the SGD classifier because its PR curve is much closer to the top-right corner, and it has a greater AUC



## Multiclass Classification

- distinguish between more than two classes
- multiclass classifiers:  
LogisticRegression, RandomForestClassifier,  
GaussianNBanNB
- strictly binary classifiers: SGDClassifier, SVC



To perform multiclass classification with multiple binary classifiers

## ① **one-versus-the-rest (OvR)** or **one-versus-all (OvA)** strategy

Example. classify the digit images into 10 classes (from 0 to 9)

- train 10 binary classifiers, one for each digit
  - a 0-detector
  - a 1-detector
  - a 2-detector
  - $\vdots$
- get the decision score from each classifier for that image
- select the class whose classifier outputs the highest score



To perform multiclass classification with multiple binary classifiers

## ② one-versus-one (OvO)

Example. classify the digit images into 10 classes (from 0 to 9)

- train a binary classifier for every pair of digits
  - one to distinguish 0s and 1s
  - another to distinguish 0s and 2s
  - another for 1s and 2s
  - $\vdots$
- For  $N$  classes, train  $\frac{N(N-1)}{2}$  classifiers
  - For the MNIST problem, this means training 45 binary classifiers
  - run the image through all 45 classifiers
  - see which class wins the most duels
- Main advantage: each classifier only needs to be trained on part of the training set containing the two classes that it must distinguish



## OvR vs. OvO

- Some algorithms (like SVM classifiers) scale poorly with the size of the training set  $\rightarrow$  **OvO**  $\rightarrow$  faster to train many classifiers on small training sets than to train few classifiers on large training sets
- For most binary classification algorithms  $\rightarrow$  **OvR**
- Note that  
scikit-learn automatically detects when you try to use a binary classification algorithm for a multiclass classification task  
and  
automatically runs OvR or OvO depending on the algorithm





## Multilabel Classification

- outputs multiple binary tags
  - face-recognition classifier: multiple faces in the same picture
  - one tag per person
- multiple classes for each instance

```
1 import numpy as np
2 from sklearn.neighbors import KNeighborsClassifier
3
4 y_train_large = (y_train >= '7')
5 y_train_odd = (y_train.astype('int8') % 2 == 1)
6 y_multilabel = np.c_[y_train_large, y_train_odd]
7
8 knn_clf = KNeighborsClassifier()
9 knn_clf.fit(X_train, y_multilabel)
10
```



## Evaluating a Multilabel Classifier

(Note that selecting the right metric really depends on the project)

- ① Measure  $F_1$  score for each individual label (or any other binary classifier metric)
  - then compute the average score
  - assumes that labels are equally important (have same weight)
- ② For a non-natively multilabel classifier (such as SVC): train one model per label
  - cannot capture the dependencies between the labels
    - to solve the above issue: organize the models in a chain
    - `ClassifierChain`

